

# NoSQL & NewSQL

Thomas Heinis

# NoSQL: Overview

- Main objective: implement distributed state
  - Different objects stored on different servers
  - Same object replicated on different servers
- Main idea: give up some of the ACID constraints to improve performance
- Simple interface:
  - Write (=Put): needs to write all replicas
  - Read (=Get): may get only one
- Eventual consistency ← Strong consistency

# NoSQL


“Not Only SQL” or “Not Relational”.

Six key features:

1. Scale horizontally “simple operations”
2. Replicate/distribute data over many servers
3. Simple call level interface (contrast w/ SQL)
4. Weaker concurrency model than ACID
5. Efficient use of distributed indexes and main memory
6. Flexible schema

# Four NoSQL Categories


## Key-Value




## Graph DB



## BigTable



## Document



# Data Model

- **Tuple** = row in a relational DB
- **Document** = nested values, extensible records (think XML or JSON)
- **Extensible record** = families of attributes have a schema, but new attributes may be added
- **Object** = like in a programming language, but without methods

# Key-value Stores

- Think “file system” more than “database”
- Consistent hashing (DHT)
- Only primary index: lookup by key
- No secondary indexes
- Transactions: single- or multi-update

# Basic Idea: Key-Value Store

Table T:

key	value
k1	v1
k2	v2
k3	v3
k4	v4

keys are sorted



- API:
  - lookup(key) → value
  - lookup(key range) → values
  - getNext → value
  - insert(key, value)
  - delete(key)
- Each row has timestamp
- Single row actions atomic  
No multi-key transactions
- No query language!

# Document Stores

- A "document" = a pointerless object = e.g. JSON = nested or not = schema-less
- In addition to KV stores, may have secondary indexes
- SimpleDB, CouchDB, MongoDB, Terrastore
- Scalability:
  - Replication (e.g. SimpleDB, CouchDB – means entire db is replicated),
  - Sharding (MongoDB);
  - Both



# Document Store (MongoDB)

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.find ({"first" : "John"})  
{  
  "_id" : ObjectId("51..."),  
  "first" : "John",  
  "last" : "Doe",  
  "age" : 39  
}
```

```
> db.user.update(  
  {"_id" : ObjectId("51...")},  
  {  
    $set: {  
      age: 40,  
      salary: 7000}  
    }  
  )
```

```
> db.user.remove(  
  "first": /^J/  
)
```

# Column Family

- Most Based on **BigTable**: Google's Distributed Storage System for Structured Data
- Data Model:
  - A big table, with column families
  - Map Reduce for querying/processing
- Examples:
  - HBase, HyperTable, Cassandra

# Column Family: Pros and Cons

- Pros:
  - Supports Semi-Structured Data
  - Naturally Indexed (columns)
  - Scalable
- Cons
  - Poor for interconnected data

# Graph Databases

- Data Model:
  - Nodes and Relationships
- Examples:
  - Neo4j, OrientDB, InfiniteGraph, AllegroGraph

# Graph Databases: Pros and Cons

- Pros:
  - Powerful data model, as general as RDBMS
  - Connected data locally indexed
  - Easy to query
- Cons
  - Sharding (lots of people working on this)
    - Scales UP reasonably well
  - Requires rewiring your brain

# Scalable Relational Systems (NewSQL)

- Means RDBMS that offer sharding
- **Key difference:**
  - NoSQL difficult or impossible to perform large-scope operations and transactions (to ensure performance)
  - Scalable RDBMS do not **preclude** these operations, but users pay a price only when they need them
- MySQL Cluster, VoltDB, Clusterix, ScaleDB, Megastore (the new BigTable)
- Many more **NewSQL** systems becoming available...

# Scalable Data Processing

- Parallel execution achieves greater efficiency
- But, parallel programming is hard
  - Parallelization
  - Fault Tolerance
  - Data Distribution
  - Load Balancing

# MapReduce (Hadoop and others)

- *“MapReduce is a programming model and an associated implementation for processing and generating large data sets”*
- Programming model
  - Abstractions to express simple computations
- Library
  - Takes care of the gory stuff: Parallelization, Fault Tolerance, Data Distribution and Load Balancing



# NoSQL Systems

- Key Value Stores

Key	Value
SW3 3TB	London
B18 4BJ	Birmingham



- Document Stores



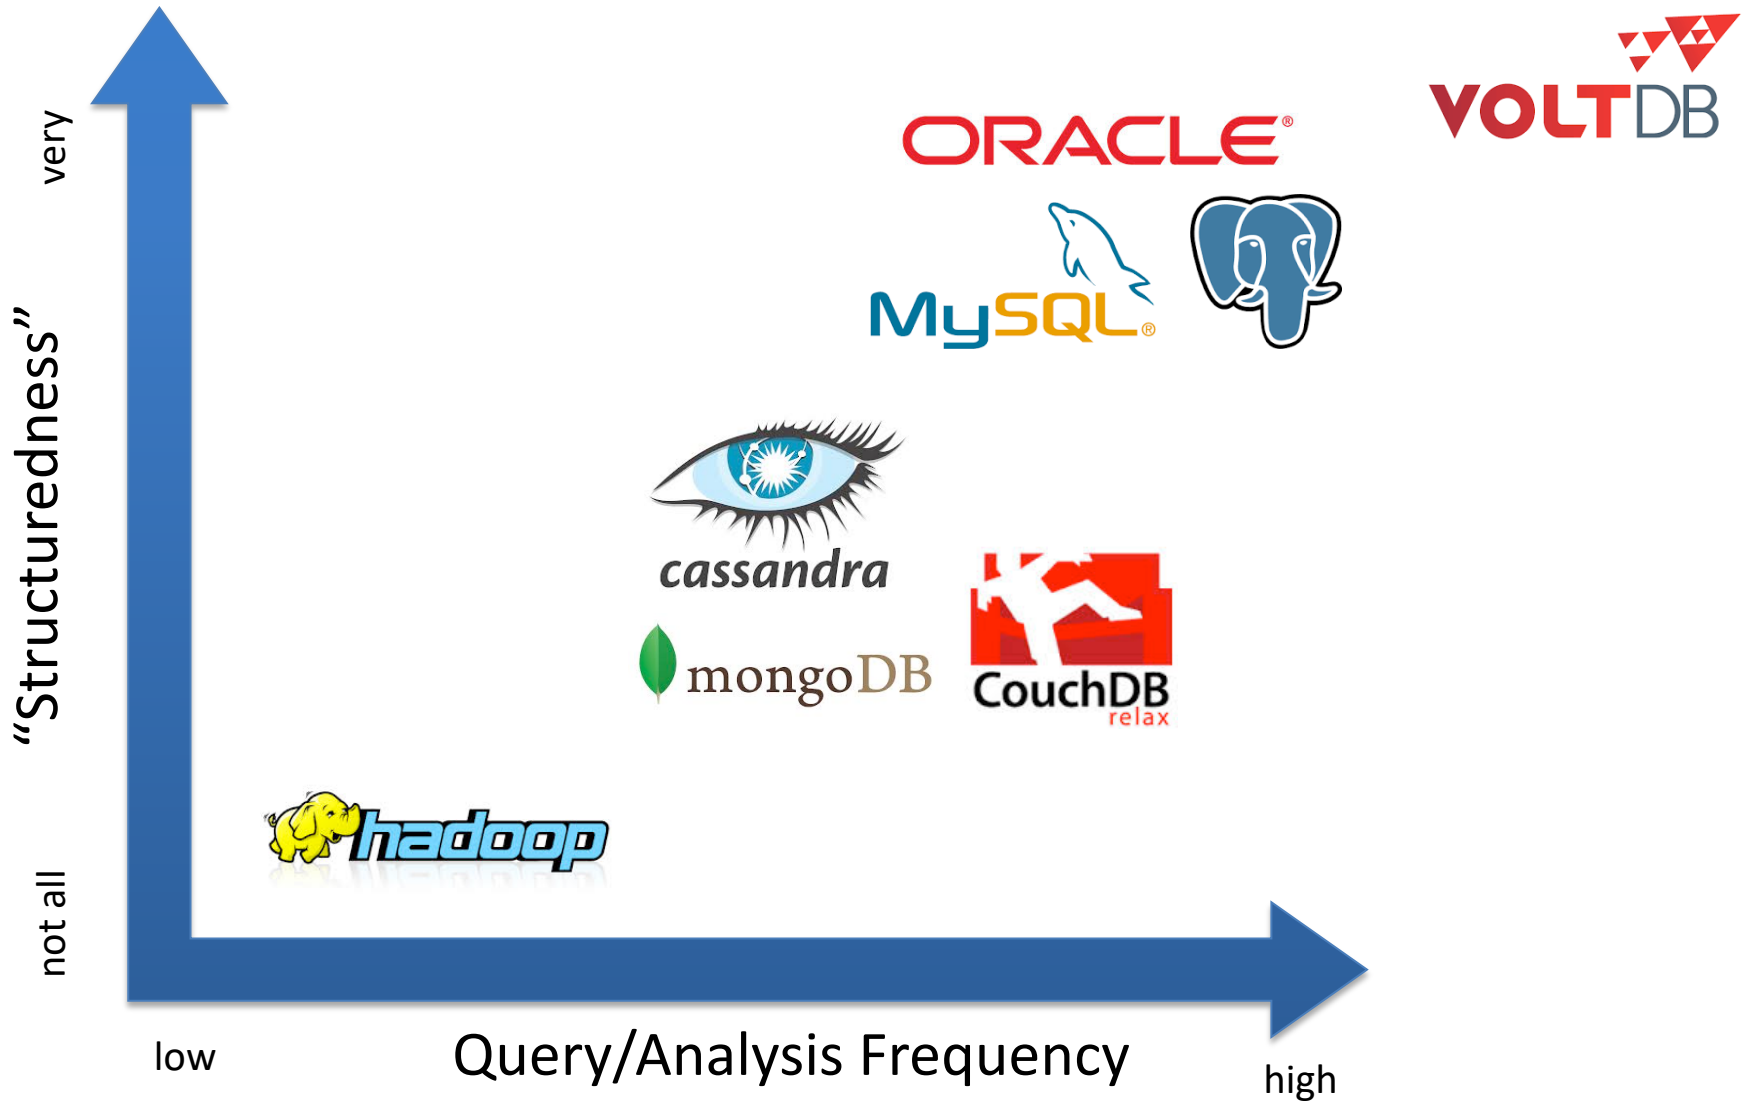
- Scalable SQL Systems



- Data Processing Systems



# Data Management Landscape\*



\*according to me