


# Lecture 2

## Our first SQL queries

Dr Fintan Nagle  
[f.nagle@imperial.ac.uk](mailto:f.nagle@imperial.ac.uk)



# Reading

Video lectures:

2.4.1 - SELECT, WHERE and LIMIT syntax.mp4

2.4.2 - Formatting conventions.mp4

3.2 - Making counting queries.mp4

# Reading

Postgres documentation on the client/server model:

<https://www.postgresql.org/docs/10/static/tutorial-arch.htm> |

Postgres documentation on aggregate functions:

<https://www.postgresql.org/docs/10/static/tutorial-agg.html>

Postgres documentation on SELECT, including GROUP BY:

<https://www.postgresql.org/docs/current/static/sql-select.htm> |

Postgres documentation on ORDER BY:

<https://www.postgresql.org/docs/current/static/queries-order.html>

NULL and the billion-dollar mistake:

<https://www.lucidchart.com/techblog/2015/08/31/the-worst-mistake-of-computer-science/>

Postgres documentation on datatypes:

<https://www.postgresql.org/docs/current/static/datatype.html>

# Vital documents

- SQL cheat sheet  
(available in the exam)

- ER diagrams for Northwind
- and dvdrental

## PostgreSQL cheat sheet

### Our practice databases

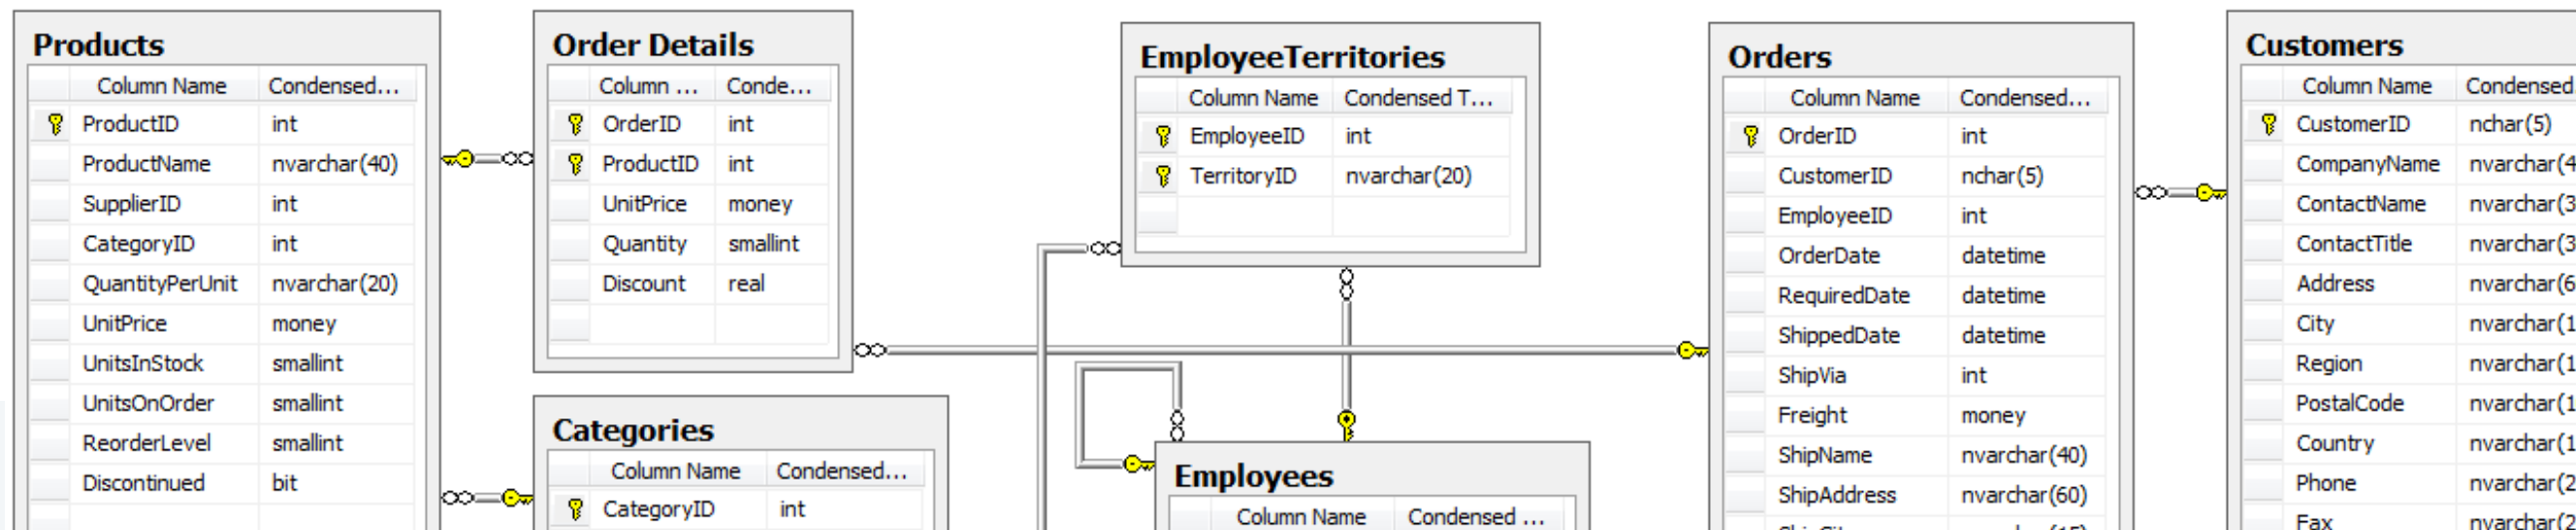
- Host: imperial.ckp3dl3vzxoh.eu-west-2.rds.amazonaws.com
- Username: imperial
- Password: imperial
- Port: 5432
- Database (maintenance DB): either **dvdrental** or **northwind**

### In Postgres

- Single quotes only - no double quotes
- This is a Postgres cheat sheet. MySQL or other versions of SQL may operate slightly differently.

### Standard query structure

- SELECT (including window functions)
- FROM
- JOINS, each with an ON
- WHERE
- ORDER BY
- LIMIT



A blue-tinted photograph of a modern glass building with a person sitting on a bench in the foreground. The building's glass facade reflects the surrounding environment. In the foreground, a person is sitting on a white, modern-style bench. The ground is paved with light-colored tiles.

# General tips for writing queries



# General SQL tips

Use newlines liberally to separate the parts of queries.

```
SELECT * FROM employees INNER JOIN roles ON employees.role_id =  
roles.id INNER JOIN managers ON employees.manager_id = managers.id  
WHERE employees.age > 50 ORDER BY employees.age DESC
```

```
SELECT * FROM  
employees INNER JOIN roles  
ON employees.role_id = roles.id  
INNER JOIN managers  
ON employees.manager_id = managers.id  
WHERE employees.age > 50  
ORDER BY employees.age DESC
```

# General programming/data science tips

- Try it out! You learn best in front of a live console window.
- Set yourself challenges!  
“I wonder if I can find all the students with more than two supervisors...”
- There are many online resources for SQL:  
Google “SQL” or “Postgres” along with your problem or error message  
Stack Overflow is very useful – if you post a clear question (that hasn’t been asked before), it’s usually answered within a day
- Make a personal cheat sheet to help you remember syntax

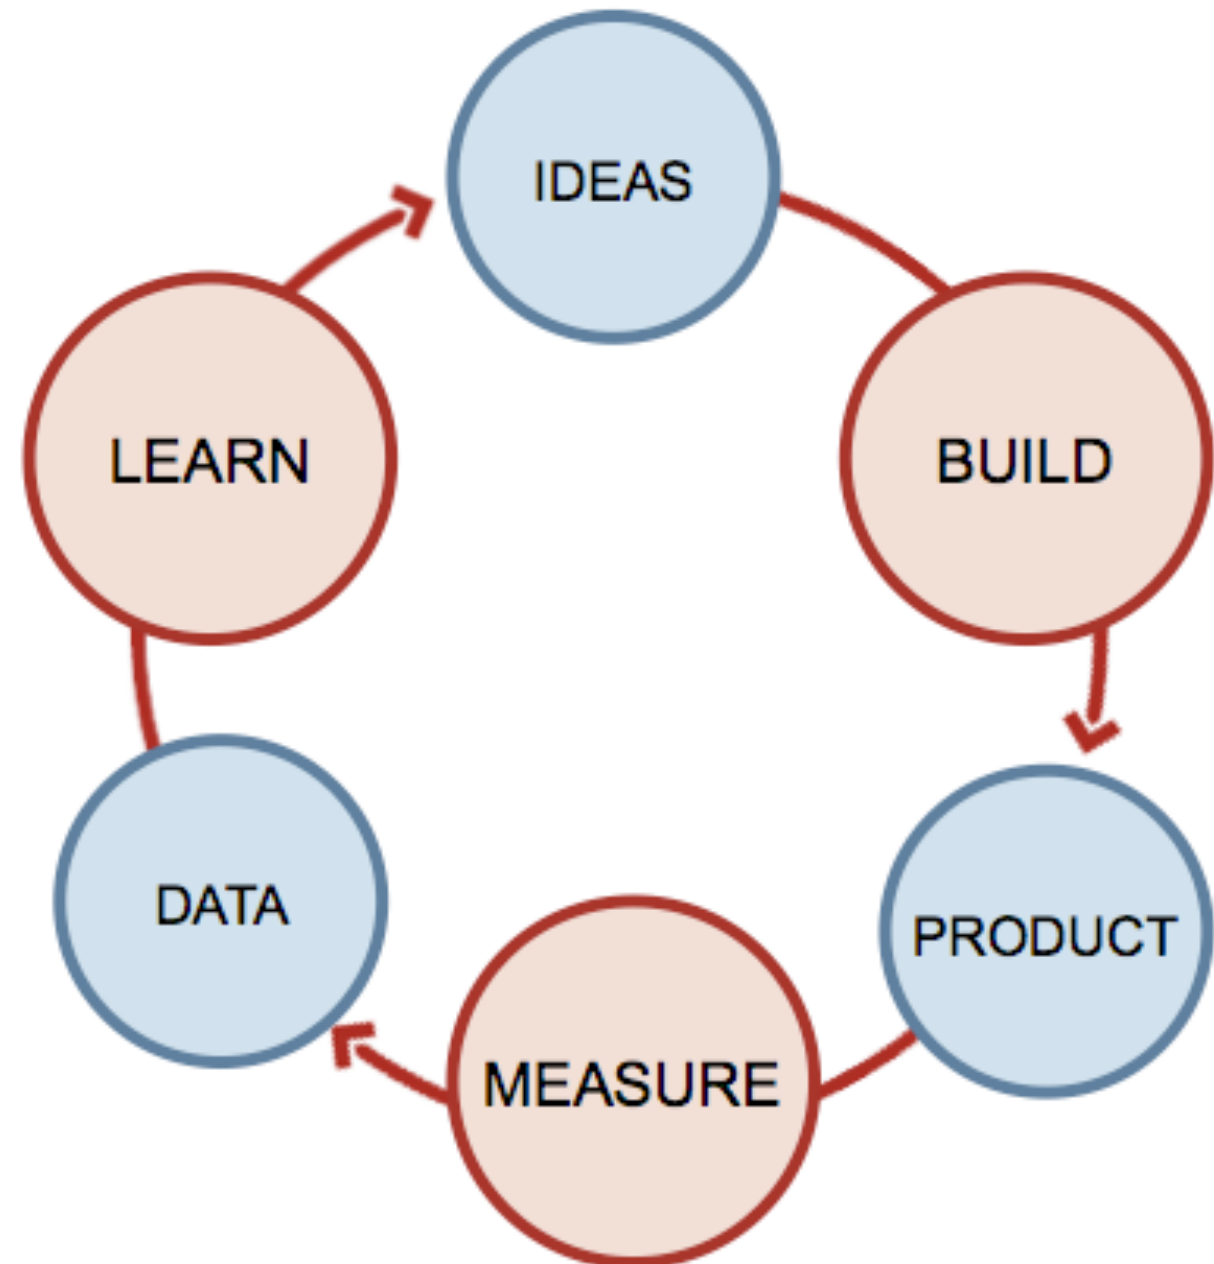
# General SQL tips

- Postgres: single quotes only, no double quotes



# Developing queries (the development loop)

- Get a simple version of the query to run first
- Check the results
- Extend and tune the query



# Naming rules

- Always use the singular (movie, not movies)
- Have good case rules:

CamelCaseMeansMoreCognitiveLoad  
snake\_case\_is\_better

# Vital documents

- SQL cheat sheet  
(available in the exam)

## PostgreSQL cheat sheet

### Our practice databases

- Host: imperial.ckp3dl3vzxoh.eu-west-2.rds.amazonaws.com
- Username: imperial
- Password: imperial
- Port: 5432
- Database (maintenance DB): either **dvdrental** or **northwind**

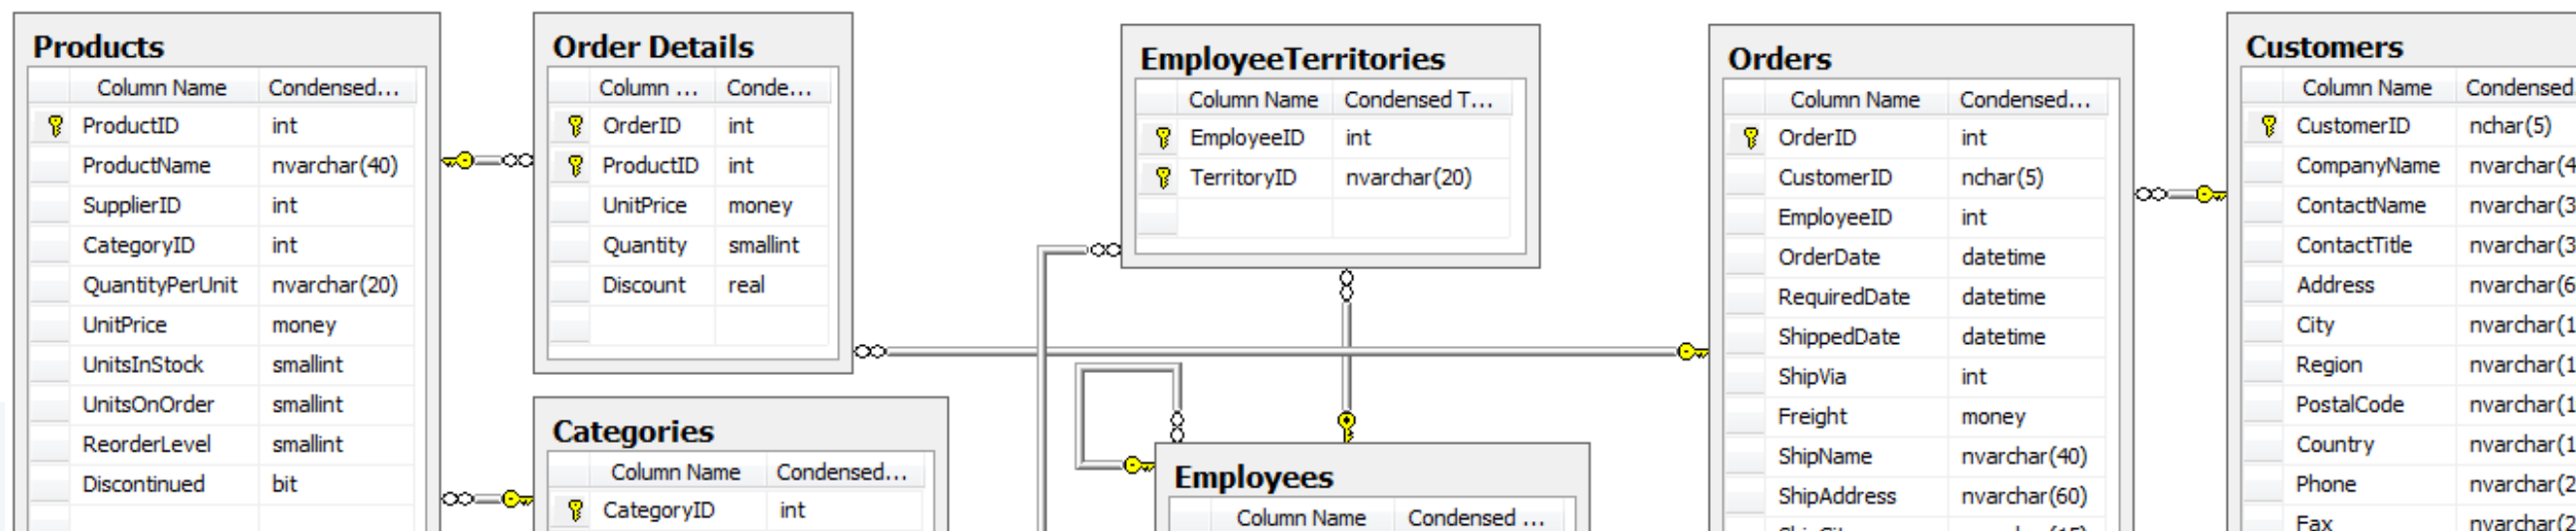
### In Postgres

- Single quotes only - no double quotes
- This is a Postgres cheat sheet. MySQL or other versions of SQL may operate slightly differently.

### Standard query structure

- SELECT (including window functions)
- FROM
- JOINS, each with an ON
- WHERE
- ORDER BY
- LIMIT

- ER diagrams for Northwind
- and dvdrental



# NULL

- NULL is a special code indicating that a value is missing, empty, or not present.
- Always in capitals in Postgres.
- When Tony Hoare introduced NULL to the Algol language in 1965, it eventually caused so many problems that he called it "my billion-dollar mistake."
- We can check for TRUE using **WHERE attribute = TRUE**  
However we can't use **WHERE attribute = NULL**  
because **nothing is ever equal to NULL**.  
So we have to use **WHERE attribute IS NULL**

# Always consider NULLs!

Guess at the result of these comparisons, and then try them out in your console. Remember that nothing is ever equal to NULL, and that NULL compared to anything is NULL.

1. **SELECT (TRUE = TRUE)**

2. **SELECT (TRUE = FALSE)**

3. **SELECT (3 < 10)**

4. **SELECT (NULL = TRUE)**

5. **SELECT (NULL = FALSE)**

6. **SELECT (NULL < 10)**

7. **SELECT (NULL <> 5)**

7. **SELECT (NULL <> 5) IS NULL**

# Always consider NULLs!

Guess at the result of these comparisons, and then try them out in your console. Remember that nothing is ever equal to NULL, and that NULL compared to anything is NULL.

1. **SELECT (TRUE = TRUE)** true

2. **SELECT (TRUE = FALSE)** false

3. **SELECT (3 < 10)** true

4. **SELECT (NULL = TRUE)** null

5. **SELECT (NULL = FALSE)** null

6. **SELECT (NULL < 10)** null


7. **SELECT (NULL <> 5)** null

7. **SELECT (NULL <> 5) IS NULL** true



# Always consider NULLs!

NULL, when included in any comparison (using = or > or < or <>), makes the result NULL.



Keys:  
uniquely identifying rows

# Keys

- Generally, a **key** is an attribute (or group of attributes) that uniquely identifies a row. For example, student ID is a key for a student table; however, DOB and last name could be a workable key too.
- Someone has to make sure that the keys are unique (we can ask Postgres to check this)

# Are rows uniquely identifiable?

A

**dogs**

name	role	age
Artemis	working	2
Fido	companion	4
Rover	companion	6
Roger	working	3
Buddy	working	7
Daisy	show	3
Shep	working	10

# Are rows uniquely identifiable?

## B

### dogs

name	role	age
Artemis	working	2
Fido	companion	4
Rover	companion	6
Roger	working	3
Buddy	working	7
Daisy	show	3
Shep	working	10
Fido	companion	4
Rover	companion	6

# Are rows uniquely identifiable?

## C

### dogs

id	name	role	age
1	Artemis	working	2
3	Fido	companion	4
4	Rover	companion	6
5	Roger	working	3
7	Buddy	working	7
8	Daisy	show	3
12	Shep	working	10



# Keys

name	breed
Rover	Great Dane
Fido	Dalmatian
Woody	Mutt
Buddy	Poodle
Jones	Great Dane

*Here name could be a key*

# Keys

The key allows UNIQUE identification of rows.

It can be a **composite key** (more than one column).

There must be a guarantee that the key will always identify rows uniquely.

This is one of the jobs of Postgres: ensuring data integrity.

name	breed
Rover	Great Dane
Fido	Dalmatian
Rover	Mutt
Buddy	Poodle
Jones	Great Dane

*Here name + breed looks like it could be a composite key*

# Keys

The key allows UNIQUE identification of rows.

It can be a **composite key** (more than one column).

There must be a guarantee that the key will always identify rows uniquely.

This is one of the jobs of Postgres: ensuring data integrity.

name	breed
Rover	Great Dane
Fido	Dalmatian
Rover	Mutt
Buddy	Poodle
Jones	Great Dane
Rover	Great Dane

*Now name + breed is obviously no good, so we need to add another key.*

# Keys

dog_id	name	breed
1	Rover	Dalmatian
2	Fido	Dalmatian
3	Daisy	Mutt
4	Daisy	Poodle

# The SELECT keyword

## Choosing columns

We can show all columns from a table:

```
SELECT * FROM dogs
```

Equivalently:

```
SELECT dogs.* FROM dogs
```

*This is useful later on when queries refer to more than one table.*

We can also request just the columns we want:

```
SELECT name, age, breed FROM dogs
```

# Renaming columns

Any columns can be renamed:

```
SELECT name AS name_of_dog FROM dogs
```

```
SELECT name AS name_of_dog, id AS dog_number FROM dogs
```

In Postgres you can actually leave out the AS; these queries give the same results as the two above:

```
SELECT name name_of_dog FROM dogs
```

```
SELECT name name_of_dog, id dog_number FROM dogs
```

*This only renames the columns in the results of your query, not in the actual database.*



# WHERE: filtering rows by certain criteria

WHERE is used to restrict rows. We will only see the rows which match our criteria.

```
SELECT * FROM employees  
WHERE age < 40
```

You can use the operators <, >, =, !=.

You can also use AND and OR (with brackets for clarity):

```
SELECT * FROM employees  
WHERE age > 40 OR (role = 'manager' AND authorised = TRUE)
```

# LIMIT: Restricting the number of rows

The LIMIT keyword restricts the number of rows.

It always comes last.

The LIMIT keyword is very useful for getting a quick look at some example data.

```
SELECT * FROM orders
```

*If the orders table is very big, this could take minutes to run, and return millions of rows!*

```
SELECT * FROM orders
```

```
LIMIT 100
```

*This will run in milliseconds. Why?*

# Inside the SQL server

The SQL processor's job is to read your query and work out how to produce your result.

This might be a big computational job!

In the end, the SQL server always has to translate your query into *disk commands* for the hard drive or SSD. This is where most of your data is stored (SQL does cache some in RAM, which is faster – but RAM is small and expensive).



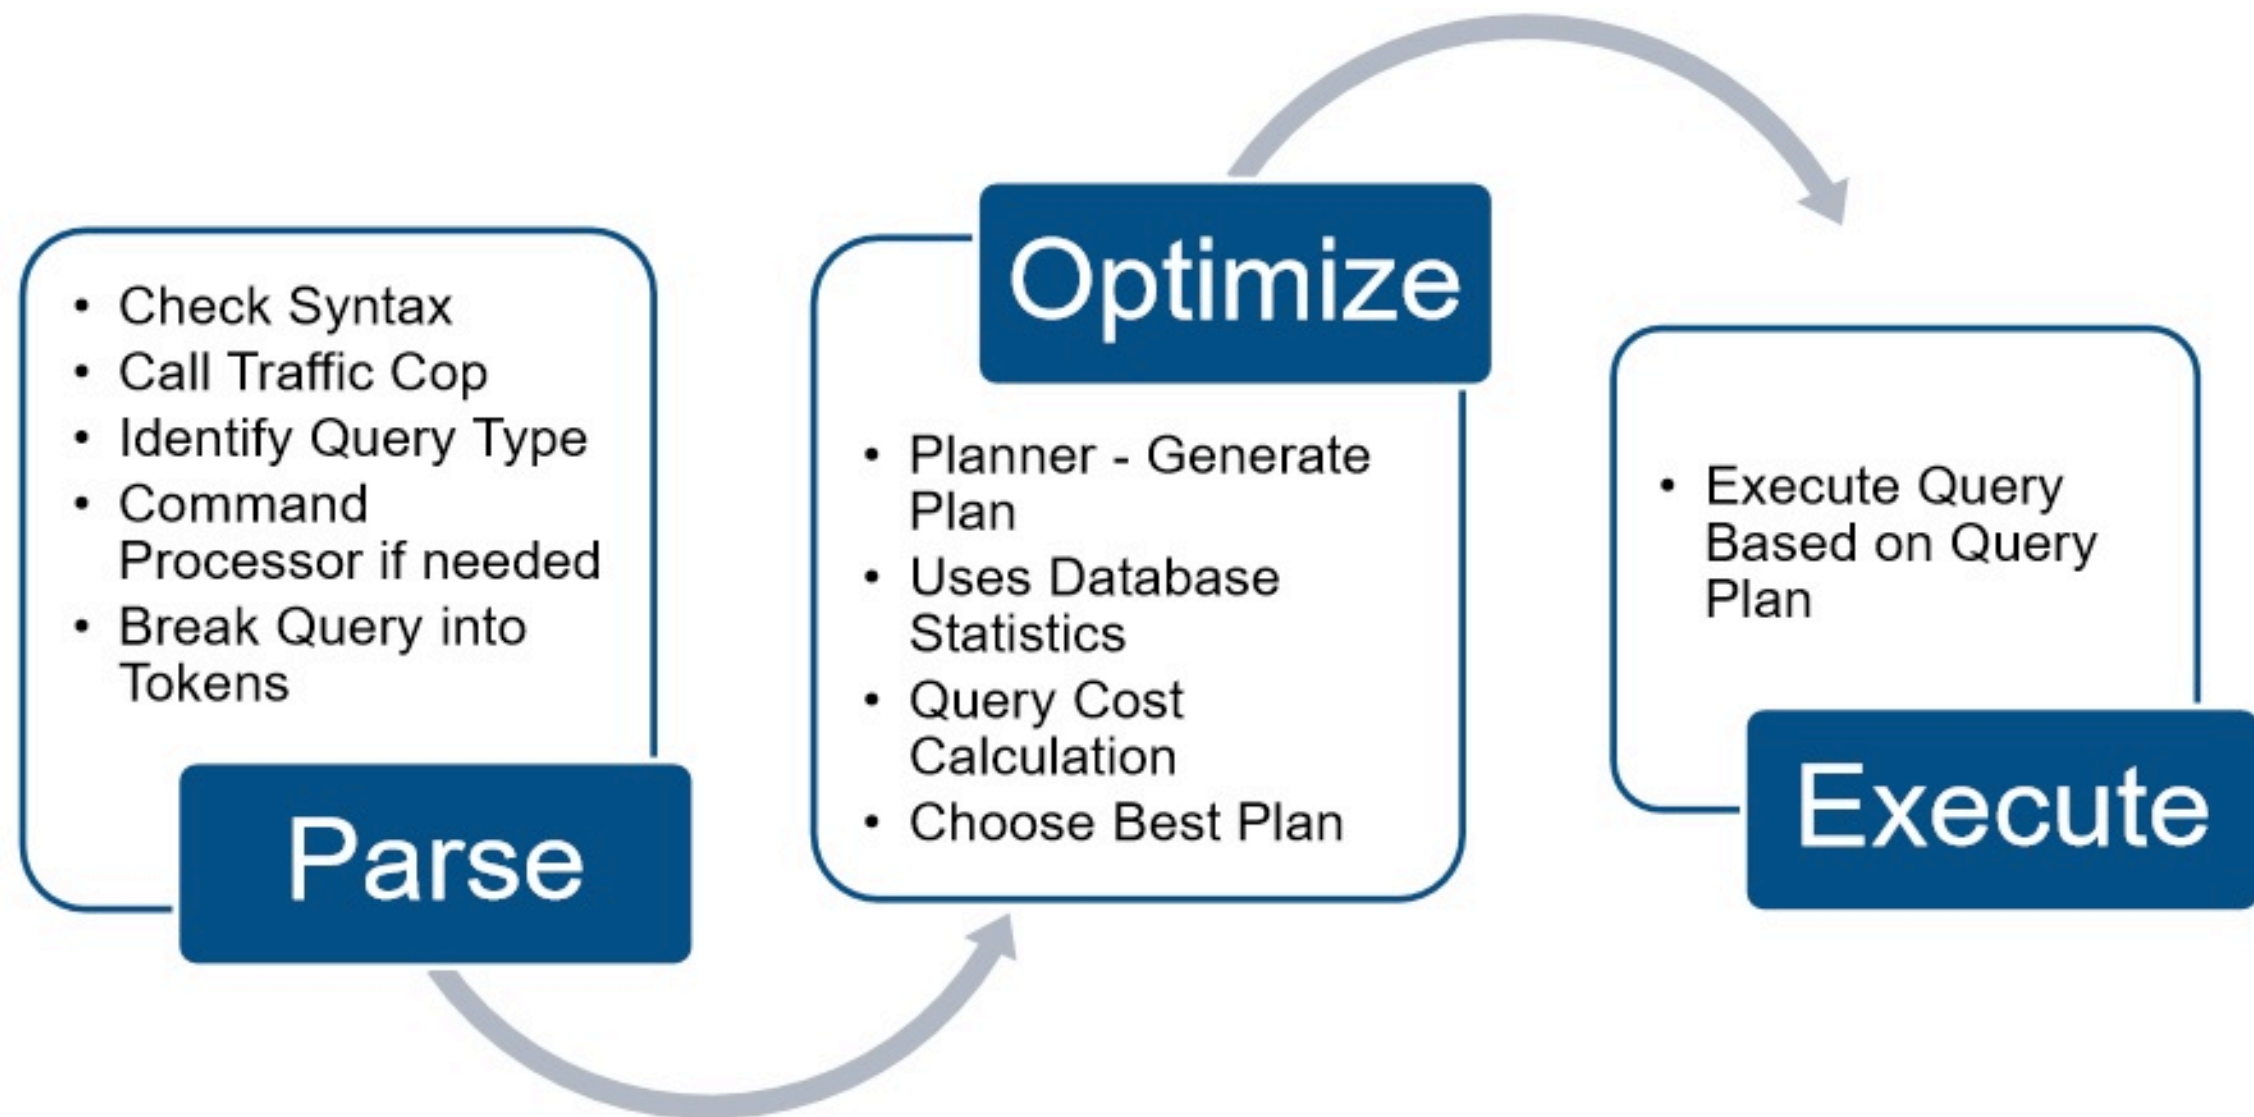
# Query optimisations

The SQL processor can make **optimisations** to save time.

For example, if your query says LIMIT 10, the processor will not have to scan the entire table, so the query will be much faster. However, if you request the MAX(age) of all employees, each row has to be checked – the whole table has to be read from disk and checked.

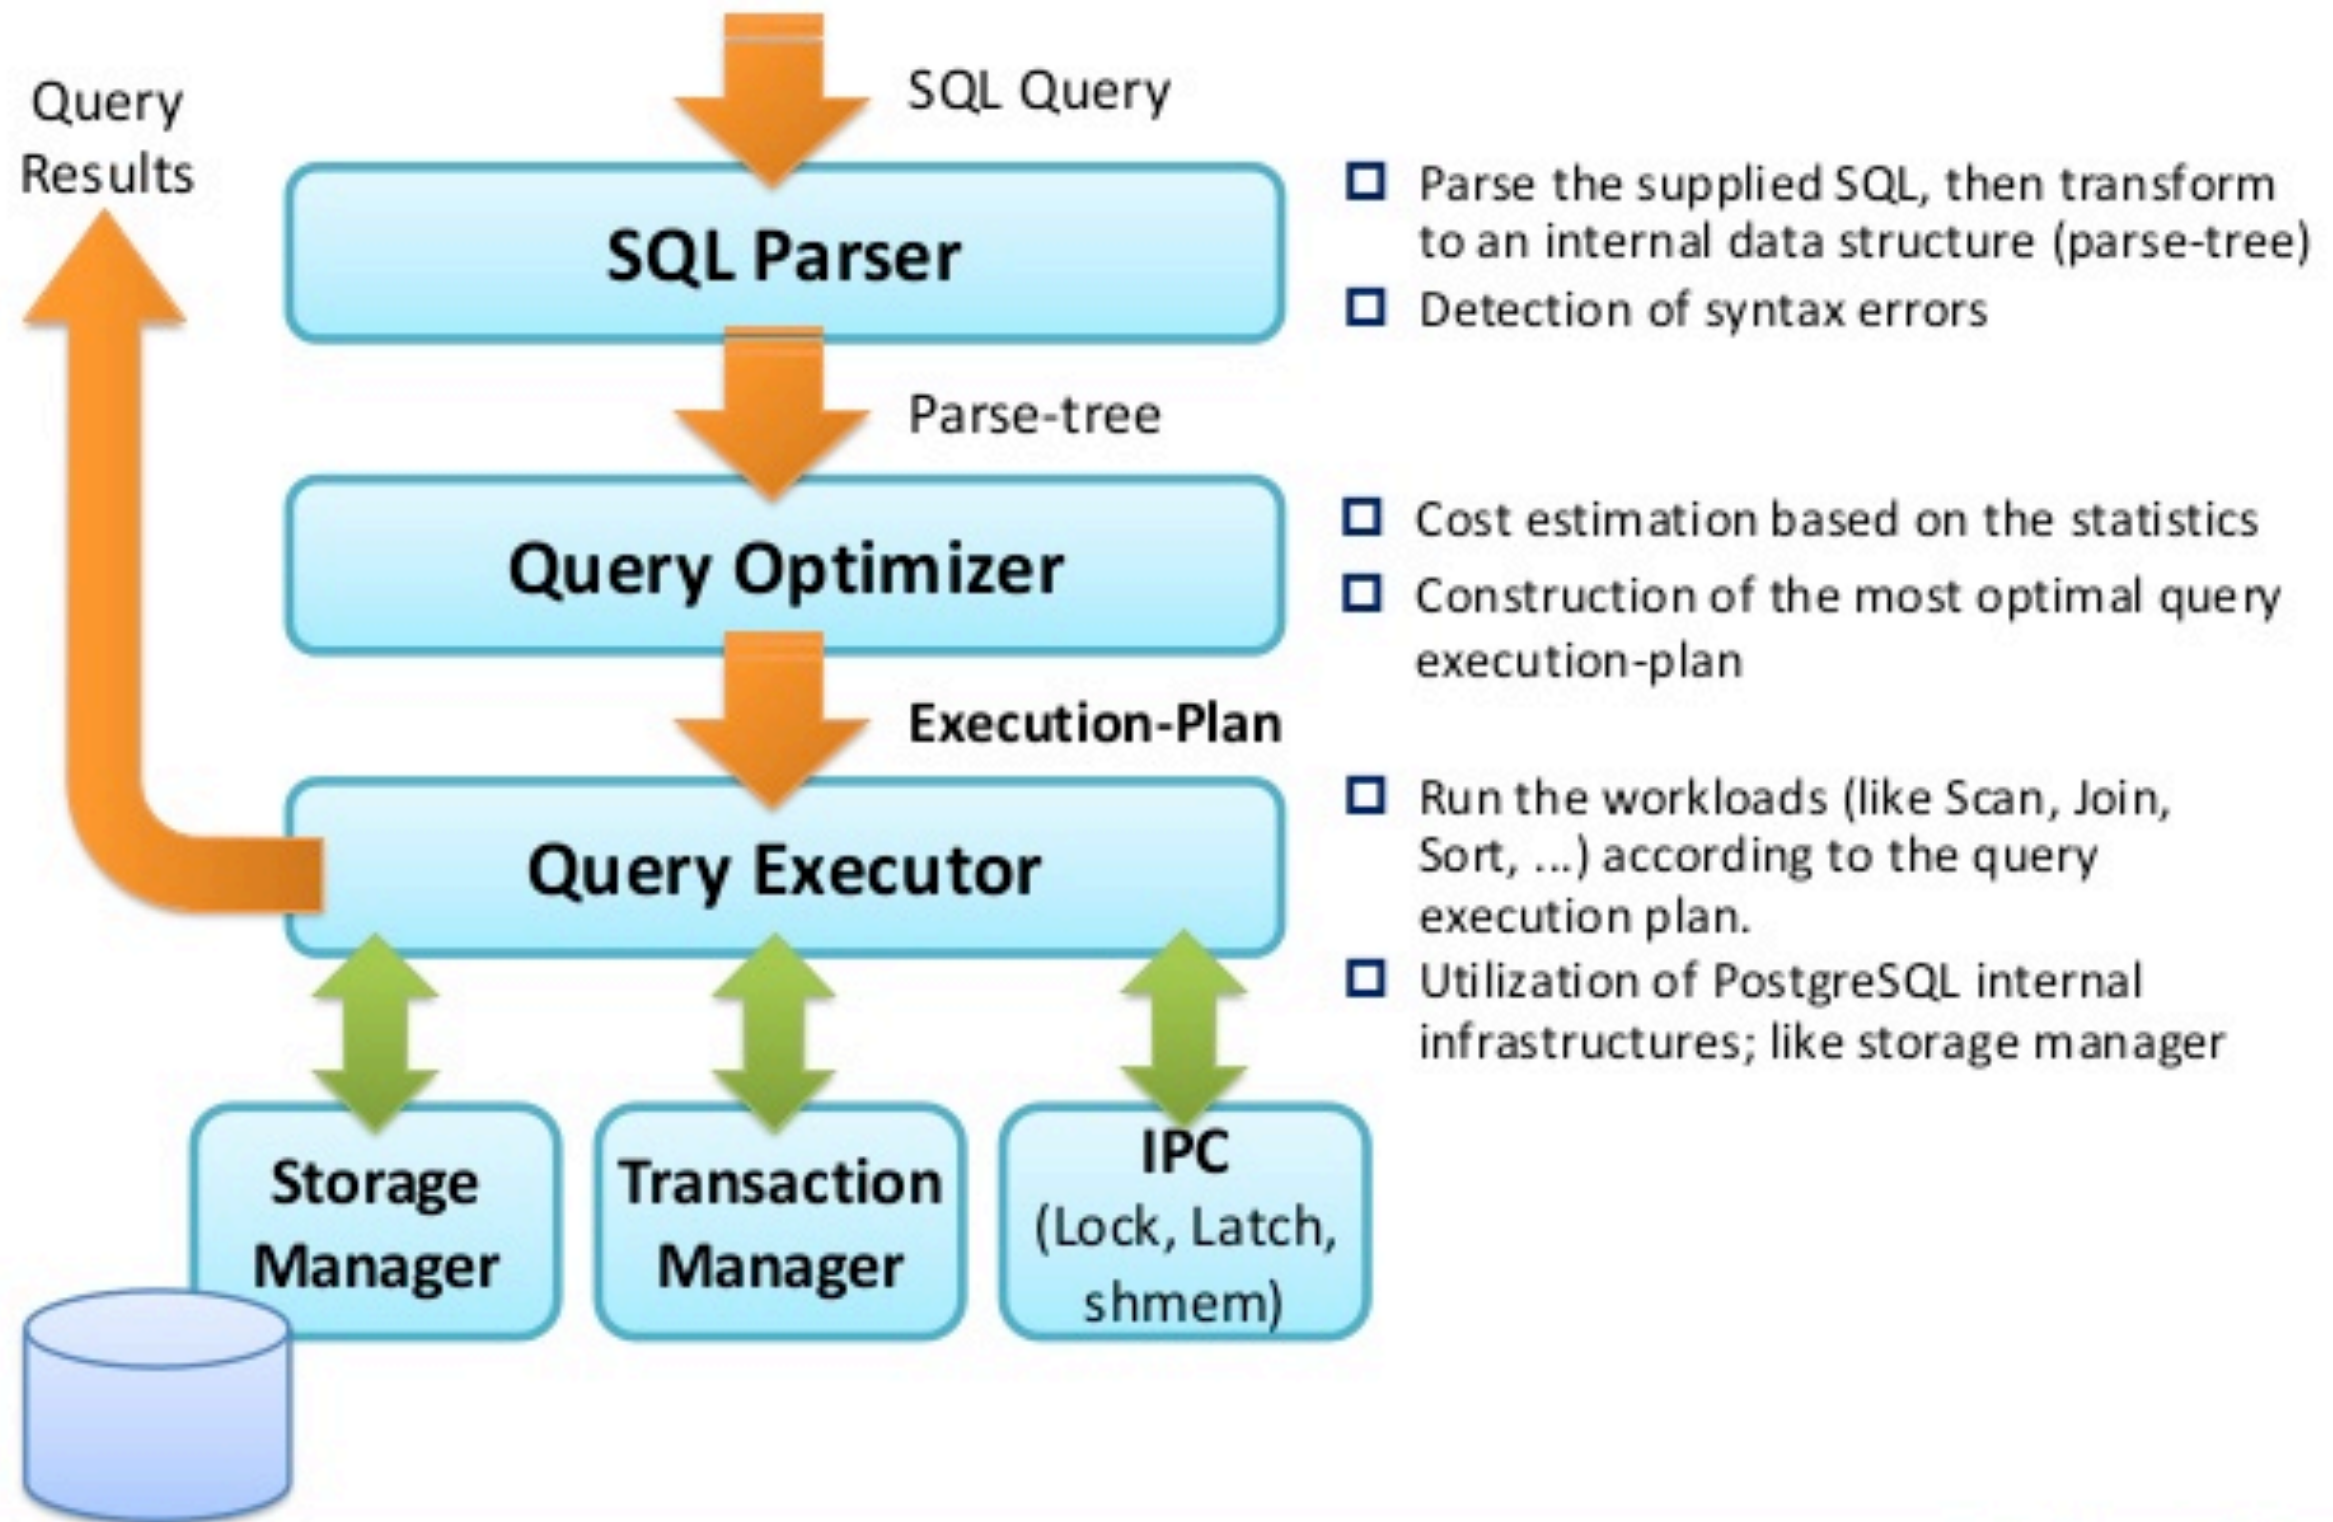
The SQL processor is not as clever as a human analyst. Later, we will see how to optimise queries ourselves.

# Inside the SQL server



("The traffic cop is the agent that is responsible for differentiating between simple and complex query commands.")

# PostgreSQL Internals







# Counting

# A note about COUNT

- **COUNT(\*)** counts the number of rows, even if they have null cells
- **COUNT(column)** counts the number of rows where that column is not null
- **COUNT(DISTINCT column)** counts the number of distinct values in that column (nulls not counted)
- Usually, you either want **COUNT(\*)** or **COUNT(DISTINCT column)**
- Whether you use **COUNT(\*)** or **COUNT(column)** can be important in non-inner joins, where rows may have missing values.

# Other aggregate functions

An aggregate function is a function which takes a group of rows, and outputs *a single number*.

- SUM
- MIN
- MAX
- AVG
- COUNT
- CORR
- STDDEV
- VARIANCE

<https://www.postgresql.org/docs/9.5/functions-aggregate.html>

# String matching

This is done with **LIKE**. **ILIKE** is the case-insensitive version.

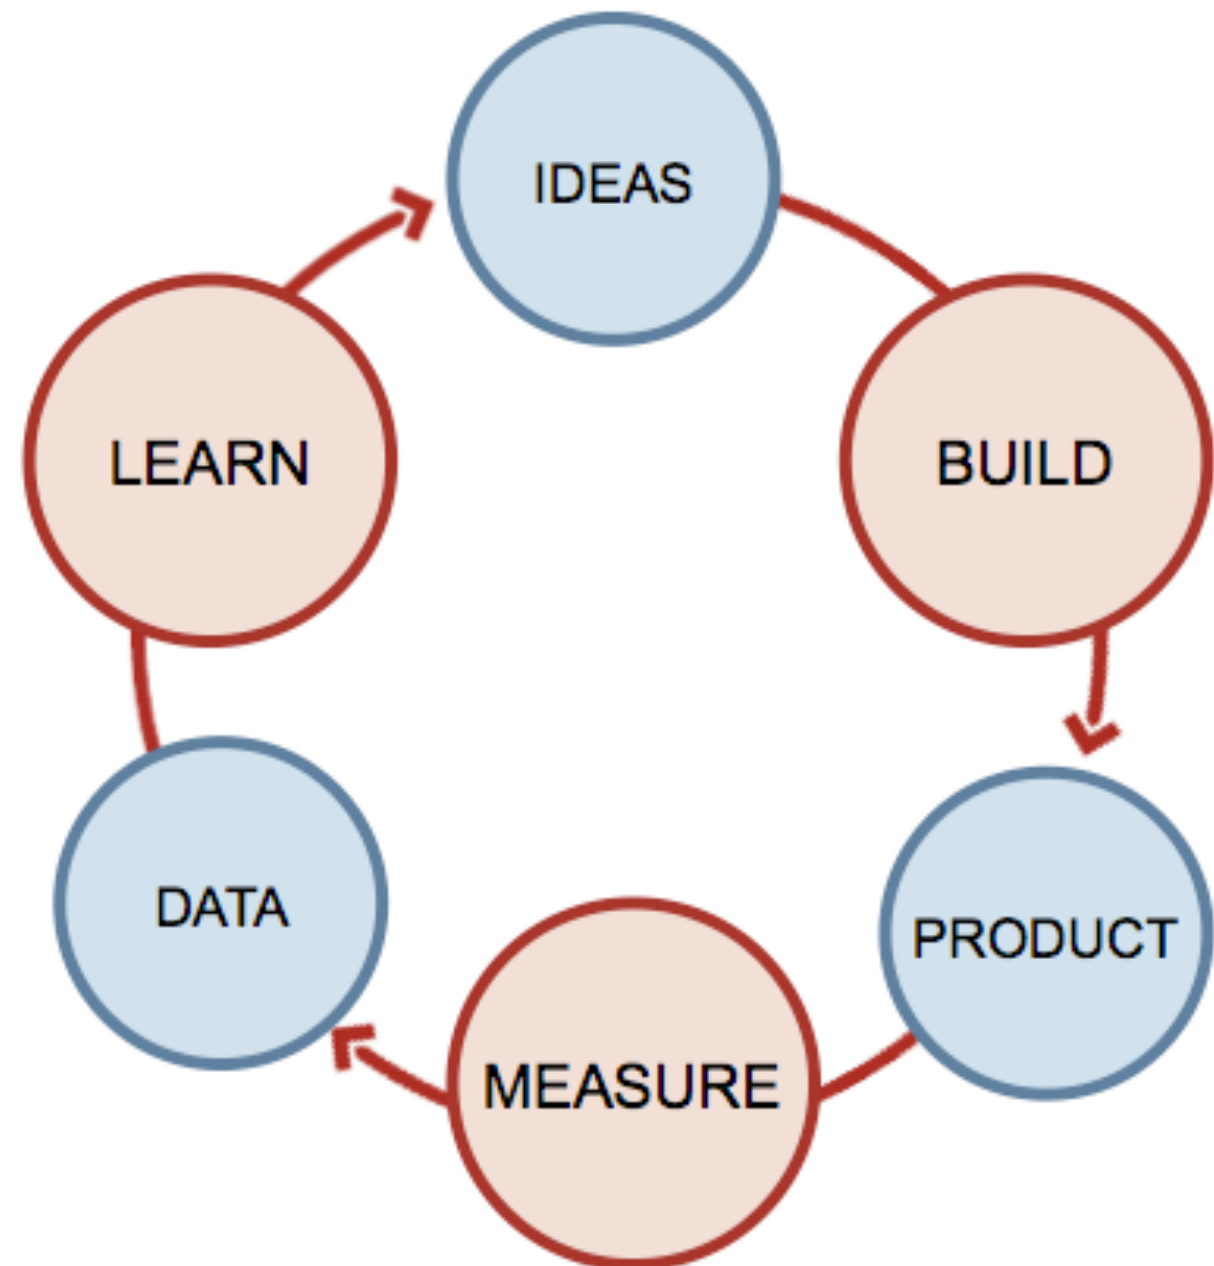
```
SELECT * FROM employees  
WHERE name LIKE 'Frank'
```

*This is just equality checking! The real power comes with from including wildcards:*

1. **WHERE** name **LIKE** '%Frank'  
*matches anything, then Frank (words ending with Frank)*
2. **WHERE** name **LIKE** 'Frank%'  
*matches Frank, then anything (names starting with Frank)*
3. **WHERE** name **LIKE** '%Frank%'  
*matches names containing Frank anywhere*

# Developing queries (the development loop)

- Get a simple version of the query to run first
- Check the results
- Extend and tune the query



# The structure of a query

A simple query:

- SELECT
- FROM
- WHERE
- ORDER BY
- LIMIT

A more complex query:

- SELECT
- FROM
- **JOINS, each with an ON**
- WHERE
- **GROUP BY**
- ORDER BY
- LIMIT